

¹ JAXtronomy: A JAX port of lenstronomy

² Alan Huang  ¹, Simon Birrer  ¹, Natalie B. Hogg  ², Aymeric Galan  ^{3,4},
³ Daniel Gilman  ⁵, Anowar J. Shajib  ^{5,6,7}, and Nan Zhang  ⁸

⁴ 1 Department of Physics and Astronomy, Stony Brook University, Stony Brook, NY 11794, USA
⁵ Laboratoire Univers et Particules de Montpellier, CNRS and Université de Montpellier (UMR-5299),
⁶ 34095 Montpellier, France ³ Max-Planck-Institut für Astrophysik, Karl-Schwarzschild Straße 1, 85748
⁷ Garching, Germany ⁴ Technical University of Munich, TUM School of Natural Sciences, Physics
⁸ Department, James-Franck-Straße 1, 85748 Garching, Germany ⁵ Department of Astronomy and
⁹ Astrophysics, University of Chicago, Chicago, Illinois 60637, USA ⁶ Kavli Institute for Cosmological
¹⁰ Physics, University of Chicago, Chicago, IL 60637, USA ⁷ Center for Astronomy, Space Science and
¹¹ Astrophysics, Independent University, Bangladesh, Dhaka 1229, Bangladesh ⁸ Department of Physics,
¹² University of Illinois, 1110 West Green St., Urbana, IL 61801, USA

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright ²³ and release the work under a ²⁴ Creative Commons Attribution 4.0 ²⁵ International License ([CC BY 4.0](#)). ²⁶ ²⁷

¹³ Summary

¹⁴ Gravitational lensing is a phenomenon where light bends around massive objects, resulting in ¹⁵ distorted images seen by an observer. Studying gravitationally lensed objects can give us key ¹⁶ insights into cosmology and astrophysics, such as constraints on the expansion rate of the ¹⁷ universe and dark matter models.

¹⁸ Thus, we introduce JAXtronomy, a re-implementation of the gravitational lensing software ¹⁹ package lenstronomy¹ ([Birrer, 2021; Birrer & Amara, 2018](#)) using JAX². JAX is a Python ²⁰ library that uses an accelerated linear algebra (XLA) compiler to improve the performance of ²¹ computing software. Our core design principle of JAXtronomy is to maintain an identical API ²² to that of lenstronomy.

²³ The main JAX features utilized in JAXtronomy are just-in-time-compilation, which can lead to ²⁴ significant reductions in execution time, and automatic differentiation, which allows for the ²⁵ implementation of gradient-based algorithms that were previously impossible. Additionally, ²⁶ JAX allows code to be run on GPUs or parallelized across CPU cores, further boosting the ²⁷ performance of JAXtronomy.

²⁸ Statement of need

²⁹ lenstronomy has been widely applied to numerous science cases, with more than 200 ³⁰ publications making use of the software, and has an increasing number of dependent packages ³¹ relying on features of lenstronomy. For instance, science cases directly involving lenstronomy ³² include galaxy evolution studies using strong lensing ([Anowar J. Shajib et al., 2021; Sheu et al., 2025; Tan et al., 2024](#)) and detailed lens modeling for measuring the Hubble constant ³³ using time-delay cosmography by the TDCOSMO collaboration ([Birrer, S. et al., 2020; Birrer, Simon & Treu, Tommaso, 2021; Collaboration et al., 2025; Gilman, D. et al., 2020; Millon, M. et al., 2020; Schmidt et al., 2025; A. J. Shajib et al., 2022; ?](#)).

³⁷ Examples of packages dependent on lenstronomy for general-purpose lensing computations and ³⁸ image modelling include the dolphin package ([Anowar J. Shajib et al., 2025](#)) for automated ³⁹ lens modeling, the galight package ([Ding et al., 2020](#)) for galaxy morphology measurements,

¹<https://github.com/lenstronomy/lenstronomy>

²<https://github.com/jax-ml/jax>

40 SLSim (Khadka et al, 2026, in prep) for simulating large populations of strong lenses, pyHalo
 41 (Gilman et al., 2019) and mejiro (Wedig et al., 2025) for simulating strong lenses with dark
 42 matter substructure, and PALTAS (Wagner-Carena et al., 2023) for neural network inference
 43 tasks.

44 In many of these applications, computational constraints are the key limiting factor for strong
 45 gravitational lensing science. For example, increased data quality and number of lenses to
 46 analyze makes lens modeling a computational bottleneck, and expensive ray-tracing through
 47 tens of thousands of dark matter substructures limit the amount of images that can be
 48 simulated, especially for the training of neural networks and simulation-based inferences. These
 49 ever-increasing computational costs have lead to the development of several JAX-accelerated
 50 and GPU-accelerated strong-lensing packages, such as gigalens (Gu et al., 2022), herculens
 51 (Galan et al., 2022), paltax (Wagner-Carena et al., 2024), GLaD (Wang et al., 2025), caustics³
 52 (Stone et al., 2024) and Google Research's jaxstronomy⁴.

53 Why JAXtronomy?

54 JAXtronomy inherits a wide range of features from lenstronomy that are not offered by any
 55 of the aforementioned JAX-accelerated or GPU-accelerated software. Such features include
 56 lenstronomy's linear amplitude solver, which reduces the number of sampled parameters
 57 during lens modeling, as well as a variety of log likelihood functions and optional punishment
 58 terms to improve robustness during fitting. JAXtronomy aims to maintain an identical API
 59 to lenstronomy so that packages dependent on lenstronomy can transition seamlessly to
 60 JAXtronomy.

61 Improvements over lenstronomy in image simulation

62 The simulation of a lensed image comes in three main steps. The first step begins with a
 63 coordinate grid in the angles seen by the observer. These coordinates are ray-traced through
 64 the deflectors back to the source plane. This process requires the calculation of light ray
 65 deflection angles at each deflector. Second, the surface brightness of the source is calculated
 66 on the ray-traced coordinate grid. This produces a lensed image. Third, the lensed image gets
 67 convolved by the point spread function (PSF) originating from diffraction of the telescope
 68 optics and atmospheric turbulence. Due to the various choices in deflector mass profiles, light
 69 model profiles, grid size, and PSF kernel size, the overall runtime of the pipeline can vary
 70 significantly.

71 In the following sections, we outline the improvements in performance that JAXtronomy has
 72 over lenstronomy for each step in the pipeline. These performance benchmarks were run using
 73 an Intel(R) Xeon(R) Gold 6338 CPU @ 2.00GHz, an NVIDIA A100 GPU, and JAX version
 74 0.7.0.

75 Deflection angle calculations

76 Each entry in the table indicates how much faster JAXtronomy is compared to lenstronomy
 77 at computing deflection angles for the corresponding deflector profile and grid size. Some
 78 comparisons vary significantly with values of function arguments, so a range is given rather
 79 than a number.

Deflector Profile	60x60 grid (cpu)	180x180 grid (cpu)	180x180 grid (gpu)
CONVERGENCE	0.4x	1.1x	0.5x
CSE	1.6x	2.6x	2.6x

³<https://github.com/Ciela-Institute/caustics>

⁴<https://github.com/google-research/google-research/tree/master/jaxstronomy>

Deflector Profile	60x60 grid (cpu)	180x180 grid (cpu)	180x180 grid (gpu)
EPL	5.1x - 15x	9.2x - 17x	37x - 120x
EPL (jax) vs	1.4x	3.0x	13x
EPL_NUMBA			
EPL_MULTIPOLE_M1M3M4	2.1x - 7x	6.4x - 13x	42x - 108x
HERNQUIST	2.0x	3.4x	5.8x
HERNQUIST_ELLIPSE_CSE	3.8x	5.4x	40x
MULTIPOLE	0.9x	1.0x	8.3x - 14x
MULTIPOLE_ELL	1.5x - 2.1x	2.0x - 2.8x	70x
NIE/SIE	0.5x	0.5x	2.0x
NFW	1.6x	3.3x	4.5x
NFW_ELLIPSE_CSE	4.1x	6.7x	31x
PJAFFE	1.0x	1.2x	2.8x
PJAFFE_ELLIPSE_POTENTIAL	1.4x	1.6x	3.1x
SHEAR	0.7x	2.0x	0.9x
SIS	1.4x	3.3x	2.0x
TNFW	2.4x	5.8x	7.5x

80 JAXtronomy computes deflection angles slower than lenstronomy when using computationally
 81 simple deflector profiles on small enough grids. This is because function call overheads are
 82 significantly higher in JAX than in standard Python, so computations that are already fast in
 83 Python can end up slower in JAX. However, these deflector profiles have higher gains from
 84 parallelization.

85 Flux calculations

86 An analogous table for the different light profiles is shown below. The MULTI_GAUSSIAN
 87 and MULTI_GAUSSIAN_ELLIPSE profiles include five GAUSSIAN and GAUSSIAN_ELLIPSE
 88 components, respectively, highlighting JAX's improved performance in sequential computations.

Light Profile	60x60 grid (cpu)	180x180 grid (cpu)	180x180 grid (gpu)
CORE_SERSIC	2.0x	6.7x	4.2x
GAUSSIAN	1.0x	2.5x	1.3x
GAUSSIAN_ELLIPSE	1.5x	3.6x	2.0x
MULTI_GAUSSIAN	3.7x	11x	7.8x
MULTI_GAUSSIAN_ELLIPSE	4.0x	13x	6.9x
SERSIC	1.0x	1.7x	3.9x
SERSIC_ELLIPSE	1.9x	5.7x	3.2x
SERSIC_ELLIPSE_Q_PHI	1.7x	5.5x	3.3x
SHAPELETS (n_max=6)	6.2x	3.4x	15x
SHAPELETS (n_max=10)	6.0x	4.5x	17x

89 FFT Convolution

90 We find that FFT convolution using JAX on CPU results in variable performance boosts or
 91 slowdowns compared to lenstronomy (which uses scipy's FFT convolution). On a 60x60
 92 grid, and kernel sizes ranging from 3 to 45, JAX on CPU ranges from being 1.1x to 2.9x faster
 93 than lenstronomy, with no obvious correlation to kernel size. On a 180x180 grid, and kernel
 94 sizes ranging from 9 to 135, JAXtronomy on CPU ranges from being 0.7x to 2.5x as fast as
 95 lenstronomy, with no obvious correlation to kernel size.

96 However, FFT convolution using JAX on GPU is significantly faster than `scipy`. On a 60x60
 97 grid, and kernel sizes ranging from 3 to 45, JAX on GPU ranges from being 1.5x to 3.5x
 98 faster than `lenstronomy`, with JAX performing better at higher kernel sizes. On a 180x180
 99 grid, and kernel sizes ranging from 9 to 135, JAXtronomy on GPU is about 10x to 20x as fast
 100 as `lenstronomy`, again with JAX performing better at higher kernel sizes.

101 Improvements over `lenstronomy` in lens modelling

102 The process of lens modelling involves finding best-fit parameters describing a lensed system from
 103 real data. In `lenstronomy`, this typically involves a Particle Swarm Optimizer (PSO) ([Kennedy & Eberhart, 1995](#)) for optimization and Monte Carlo Markov Chains for posterior sampling.
 104 `JAXtronomy` retains these lens modelling algorithms from `lenstronomy` while benefitting from
 105 the increased performance outlined above.

106 In the following table, we compare `JAXtronomy`'s PSO performance to that of `lenstronomy`
 107 when modeling a lens with a singular isothermal ellipsoid (SIE) mass profile, Sersic-ellipse
 108 source and lens light profile, and a quadruply-imaged point source. The image is simulated
 109 using a 100x100 grid and FFT convolved using a PSF kernel with a size of 13 pixels. These
 110 benchmarks were performed using the same hardware as in the previous section.

Device	64 Particles	128 Particles	256 Particles	512 Particles
lenstronomy (baseline)	59s	138s	245s	555s
1 CPU core	3x	3x	3x	4x
2 CPU cores	5x	6x	6x	7x
4 CPU cores	8x	12x	11x	12x
8 CPU cores	11x	17x	17x	24x
16 CPU cores	13x	20x	22x	29x
32 CPU cores	13x	20x	20x	29x
GPU	8x	7x	27x	46x

111 Additionally, using JAX's autodifferentiation, we have implemented the L-BFGS gradient
 112 descent algorithm from the `Optax`⁵ library ([DeepMind et al., 2020](#)) for optimization. This is
 113 a significant improvement over `lenstronomy`'s PSO, which does not have access to gradient
 114 information. Due to the stochastic nature of the PSO, we do not present a concrete comparison
 115 between `lenstronomy`'s PSO and `JAXtronomy`'s minimizer for the time it takes to find best-fit
 116 parameters.

118 References

- 119 Birrer, S. (2021). Gravitational lensing formalism in a curved arc basis: A continuous
 120 description of observables and degeneracies from the weak to the strong lensing regime.
 121 *The Astrophysical Journal*, 919(1), 38. <https://doi.org/10.3847/1538-4357/ac1108>
- 122 Birrer, S., & Amara, A. (2018). Lenstronomy: Multi-purpose gravitational lens modelling
 123 software package. *Physics of the Dark Universe*, 22, 189–201. <https://doi.org/10.1016/j.dark.2018.11.002>
- 125 Birrer, Simon, & Treu, Tommaso. (2021). TDCOSMO - v. Strategies for precise and accurate
 126 measurements of the hubble constant with strong lensing. *Astronomy & Astrophysics*, 649,
 127 A61. <https://doi.org/10.1051/0004-6361/202039179>
- 128 Birrer, S., Shajib, A. J., Galan, A., Millon, M., Treu, T., Agnello, A., Auger, M., Chen, G.
 129 C.-F., Christensen, L., Collett, T., Courbin, F., Fassnacht, C. D., Koopmans, L. V. E.,

⁵<https://github.com/google-deepmind/optax>

- 130 Marshall, P. J., Park, J.-W., Rusu, C. E., Sluse, D., Spinello, C., Suyu, S. H., ... Van de
 131 Vyvere, L. (2020). TDCOSMO - IV. Hierarchical time-delay cosmography – joint inference
 132 of the hubble constant and galaxy density profiles*. *Astronomy & Astrophysics*, 643, A165.
 133 <https://doi.org/10.1051/0004-6361/202038861>
- 134 Collaboration, T., Birrer, S., Buckley-Geer, E. J., Cappellari, M., Courbin, F., Dux, F.,
 135 Fassnacht, C. D., Frieman, J. A., Galan, A., Gilman, D., Huang, X.-Y., Knobel, S.,
 136 Langeroodi, D., Lin, H., Millon, M., Morishita, T., Motta, V., Mozumdar, P., Paic, E., ...
 137 Wong, K. C. (2025). *TDCOSMO 2025: Cosmological constraints from strong lensing time*
 138 *delays.* <https://arxiv.org/abs/2506.03023>
- 139 DeepMind, Babuschkin, I., Baumli, K., Bell, A., Bhupatiraju, S., Bruce, J., Buchlovsky, P.,
 140 Budden, D., Cai, T., Clark, A., Danihelka, I., Dedieu, A., Fantacci, C., Godwin, J., Jones,
 141 C., Hemsley, R., Hennigan, T., Hessel, M., Hou, S., ... Viola, F. (2020). *The DeepMind*
 142 *JAX Ecosystem.* <http://github.com/google-deepmind>
- 143 Ding, X., Silverman, J., Treu, T., Schulze, A., Schramm, M., Birrer, S., Park, D., Jahnke, K.,
 144 Bennert, V. N., Kartaltepe, J. S., Koekemoer, A. M., Malkan, M. A., & Sanders, D. (2020).
 145 The Mass Relations between Supermassive Black Holes and Their Host Galaxies at $1 < z$
 146 < 2 HST-WFC3. *The Astrophysical Journal*, 888(1), 37. <https://doi.org/10.3847/1538-4357/ab5b90>
- 148 Galan, A., Vernardos, G., Peel, A., Courbin, F., & Starck, J.-L. (2022). Using wavelets
 149 to capture deviations from smoothness in galaxy-scale strong lenses. *Astronomy &*
 150 *Astrophysics*, 668, A155. <https://doi.org/10.1051/0004-6361/202244464>
- 151 Gilman, D., Birrer, S., Nierenberg, A., Treu, T., Du, X., & Benson, A. (2019). Warm dark
 152 matter chills out: Constraints on the halo mass function and the free-streaming length of
 153 dark matter with eight quadruple-image strong gravitational lenses. *Monthly Notices of the*
 154 *Royal Astronomical Society*, 491(4), 6077–6101. <https://doi.org/10.1093/mnras/stz3480>
- 155 Gilman, D., Birrer, S., & Treu, T. (2020). TDCOSMO - III. Dark matter substructure meets
 156 dark energy. The effects of (sub)halos on strong-lensing measurements of H_0 . *Astronomy*
 157 & *Astrophysics*, 642, A194. <https://doi.org/10.1051/0004-6361/202038829>
- 158 Gu, A., Huang, X., Sheu, W., Aldering, G., Bolton, A. S., Boone, K., Dey, A., Filipp, A.,
 159 Jullo, E., Perlmutter, S., Rubin, D., Schlafly, E. F., Schlegel, D. J., Shu, Y., & Suyu, S. H.
 160 (2022). GIGA-lens: Fast bayesian inference for strong gravitational lens modeling. *The*
 161 *Astrophysical Journal*, 935(1), 49. <https://doi.org/10.3847/1538-4357/ac6de4>
- 162 Kennedy, J., & Eberhart, R. (1995). Particle swarm optimization. *Proceedings of ICNN'95*
 163 - *International Conference on Neural Networks*, 4, 1942–1948 vol.4. <https://doi.org/10.1109/ICNN.1995.488968>
- 165 Millon, M., Galan, A., Courbin, F., Treu, T., Suyu, S. H., Ding, X., Birrer, S., Chen, G. C.-F.,
 166 Shajib, A. J., Sluse, D., Wong, K. C., Agnello, A., Auger, M. W., Buckley-Geer, E. J.,
 167 Chan, J. H. H., Collett, T., Fassnacht, C. D., Hilbert, S., Koopmans, L. V. E., ... Van
 168 de Vyvere, L. (2020). TDCOSMO - i. An exploration of systematic uncertainties in the
 169 inference of H_0 from time-delay cosmography. *Astronomy & Astrophysics*, 639, A101.
 170 <https://doi.org/10.1051/0004-6361/201937351>
- 171 Schmidt, T., Treu, T., Birrer, S., Millon, M., Sluse, D., Galan, A., Shajib, A., Lemon, C.,
 172 Dux, F., & Courbin, F. (2025). TDCOSMO. XVIII. Strong lens model and time-delay
 173 predictions for J1721+8842, the first einstein zigzag lens. *Astronomy & Astrophysics*.
 174 <https://doi.org/10.1051/0004-6361/202449984>
- 175 Shajib, Anowar J., Nihal, N. S., Tan, C. Y., Sahu, V., Birrer, S., Treu, T., & Frieman, J. (2025).
 176 *dolphin: A fully automated forward modeling pipeline powered by artificial intelligence for*
 177 *galaxy-scale strong lenses.* <https://arxiv.org/abs/2503.22657>
- 178 Shajib, Anowar J., Treu, T., Birrer, S., & Sonnenfeld, A. (2021). Dark matter haloes of

- 179 massive elliptical galaxies at $z \sim 0.2$ are well described by the Navarro-Frenk-White
180 profile. *Monthly Notices of the Royal Astronomical Society*, 503(2), 2380–2405. <https://doi.org/10.1093/mnras/stab536>
- 182 Shajib, A. J., Wong, K. C., Birrer, S., Suyu, S. H., Treu, T., Buckley-Geer, E. J., Lin, H., Rusu, C.
183 E., Poh, J., Palmese, A., Agnello, A., Auger-Williams, M. W., Galan, A., Schuldt, S., Sluse,
184 D., Courbin, F., Frieman, J., & Millon, M. (2022). TDCOSMO. IX. Systematic comparison
185 between lens modelling software programs: Time-delay prediction for WGD 2038–4008.
186 *Astronomy & Astrophysics*, 667, A123. <https://doi.org/10.1051/0004-6361/202243401>
- 187 Sheu, W., Shajib, A. J., Treu, T., Sonnenfeld, A., Birrer, S., Cappellari, M., Oldham, L. J., &
188 Tan, C. Y. (2025). Project Dinos II: redshift evolution of dark and luminous matter density
189 profiles in strong-lensing elliptical galaxies across $0.1 < z < 0.9$. *Monthly Notices of the*
190 *Royal Astronomical Society*, 541(1), 1–27. <https://doi.org/10.1093/mnras/staf976>
- 191 Stone, C., Adam, A., Coogan, A., Yantovski-Barth, M. J., Philipp, A., Setiawan, L., Core,
192 C., Legin, R., Wilson, C., Barco, G. M., Hezaveh, Y., & Perreault-Levasseur, L. (2024).
193 Caustics: A python package for accelerated strong gravitational lensing simulations. *Journal*
194 *of Open Source Software*, 9(103), 7081. <https://doi.org/10.21105/joss.07081>
- 195 Tan, C. Y., Shajib, A. J., Birrer, S., Sonnenfeld, A., Treu, T., Wells, P., Williams, D.,
196 Buckley-Geer, E. J., Drlica-Wagner, A., & Frieman, J. (2024). Project Dinos I: A joint
197 lensing-dynamics constraint on the deviation from the power law in the mass profile of
198 massive ellipticals. *Monthly Notices of the Royal Astronomical Society*, 530(2), 1474–1505.
199 <https://doi.org/10.1093/mnras/stae884>
- 200 Wagner-Carena, S., Aalbers, J., Birrer, S., Nadler, E. O., Darragh-Ford, E., Marshall, P. J., &
201 Wechsler, R. H. (2023). From images to dark matter: End-to-end inference of substructure
202 from hundreds of strong gravitational lenses. *The Astrophysical Journal*, 942(2), 75.
203 <https://doi.org/10.3847/1538-4357/aca525>
- 204 Wagner-Carena, S., Lee, J., Pennington, J., Aalbers, J., Birrer, S., & Wechsler, R. H.
205 (2024). A strong gravitational lens is worth a thousand dark matter halos: Inference on
206 small-scale structure using sequential methods. *The Astrophysical Journal*, 975(2), 297.
207 <https://doi.org/10.3847/1538-4357/ad6e70>
- 208 Wang, H., Suyu, S. H., Galan, A., Halkola, A., Cappellari, M., Shajib, A. J., & Cernetic, M.
209 (2025). GPU-accelerated gravitational lensing & dynamical (GLaD) modeling for cosmology
210 and galaxies. <https://arxiv.org/abs/2504.01302>
- 211 Wedig, B., Daylan, T., Birrer, S., Cyr-Racine, F.-Y., Dvorkin, C., Finkbeiner, D. P., Huang, A.,
212 Huang, X., Karthik, R., Khadka, N., Natarajan, P., Nierenberg, A. M., Peter, A. H. G., Pierel,
213 J. D. R., Tang, X. T., & Wechsler, R. H. (2025). The roman view of strong gravitational
214 lenses. *The Astrophysical Journal*, 986(1), 42. <https://doi.org/10.3847/1538-4357/adc24f>