The Journal of Open Source Software

JAXtronomy: A JAX port of lenstronomy

Alan Huang[1], **Simon Birrer**[1], **Natalie B. Hogg**[2], **Aymeric Galan**[3,4], **Daniel Gilman**[5], **Anowar J. Shajib**[5,6,7], and **Nan Zhang**[8]

1 Department of Physics and Astronomy, Stony Brook University, Stony Brook, NY 1794, USA **2** Laboratoire Univers et Particules de Montpellier, CNRS and Université de Montpellier (UMR-5299), 34095 Montpellier, France **3** Max-Planck-Institut für Astrophysik, Karl-Schwarzschild Straße 1, 85748 Garching, Germany **4** Technical University of Munich, TUM School of Natural Sciences, Physics Department, James-Franck-Straße 1, 85748 Garching, Germany **5** Department of Astronomy and Astrophysics, University of Chicago, Chicago, Illinois 60637, USA **6** Kavli Institute for Cosmological Physics, University of Chicago, Chicago, IL 60637, USA **7** Center for Astronomy, Space Science and Astrophysics, Independent University, Bangladesh, Dhaka 1229, Bangladesh **8** Department of Physics, University of Illinois, 1110 West Green St., Urbana, IL 61801, USA

DOI: 10.xxxxxx/draft

Software

- Review
- Repository
- Archive

Editor: Open Journals

Reviewers:

- @openjournals

Submitted: 01 January 1970
Published: unpublished

Summary

`JAXtronomy` is a re-implementation of the gravitational lensing software package `lenstronomy`[1] (Birrer, 2021; Birrer & Amara, 2018) using `JAX`[2], a Python library that uses an accelerated linear algebra (XLA) compiler to improve the performance of computing software. Our core design principle of `JAXtronomy` is to maintain an identical API to that of `lenstronomy`.

The main `JAX` features utilized in `JAXtronomy` are just-in-time-compilation, which can lead to significant reductions in execution time, and automatic differentiation, which allows for the implementation of gradient-based algorithms that were previously impossible. Additionally, `JAX` allows code to be run on GPUs, further boosting the performance of `JAXtronomy`.

Statement of need

`lenstronomy` has been widely applied to numerous science cases, with more than 200 publications making use of the software, and has an increasing number of dependent packages relying on features of `lenstronomy`. For instance, science cases directly involving `lenstronomy` include galaxy evolution studies using strong lensing (Anowar J. Shajib et al., 2021; Sheu et al., 2025; Tan et al., 2024) and detailed lens modeling for measuring the Hubble constant using time-delay cosmography by the TDCOSMO collaboration (Birrer, S. et al., 2020; Birrer, Simon & Treu, Tommaso, 2021; Collaboration et al., 2025; Gilman, D. et al., 2020; Millon, M. et al., 2020; Schmidt et al., 2025; A. J. Shajib et al., 2022; **?**).

Examples of packages dependent on `lenstronomy` for general-purpose lensing computations and image modelling include the `dolphin` package (Anowar J. Shajib et al., 2025) for automated lens modeling, the `galight` package (Ding et al., 2020) for galaxy morphology measurements, `SLSim` (Khadka et al, 2025, in prep) for simulating large populations of strong lenses, `pyHalo` (Gilman et al., 2019) and `mejiro` (Wedig et al., 2025) for simulating strong lenses with dark matter substructure, and `PALTAS` (Wagner-Carena et al., 2023) for neural network inference tasks.

In many of these applications, computational constraints are the key limiting factor for strong gravitational lensing science. For example, increased data quality and number of lenses to

[1]https://github.com/lenstronomy/lenstronomy
[2]https://github.com/jax-ml/jax

analyze makes lens modeling a computational bottleneck, and expensive ray-tracing through tens of thousands of dark matter substructures limit the amount of images that can be simulated, especially for the training of neural networks and simulation-based inferences.

These ever-increasing computational costs have lead to the development of several JAX-accelerated strong-lensing packages, such as `gigalens` (Gu et al., 2022), `herculens` (Galan et al., 2022), `paltax` (Wagner-Carena et al., 2024), `GLaD` (Wang et al., 2025), and Google Research's `jaxstronomy`[3]. Such packages have been directly inspired by `lenstronomy` and/or support specific use cases. With `JAXtronomy`, we aim to support a wide range of features offered by `lenstronomy` while maintaining an identical API so that packages dependent on `lenstronomy` can transition seamlessly to `JAXtronomy`.

Improvements over lenstronomy in image simulation

The simulation of a lensed image comes in three main steps. The first step begins with a coordinate grid in the angles seen by the observer. These coordinates are ray-traced through the deflectors back to the source plane. This process requires the calculation of light ray deflection angles at each deflector. Second, the surface brightness of the source is calculated on the ray-traced coordinate grid. This produces a lensed image. Third, the lensed image gets convolved by the point spread function (PSF) originating from diffraction of the telescope optics and atmospheric turbulence. Due to the various choices in deflector mass profiles, light model profiles, grid size, and PSF kernel size, the overall runtime of the pipeline can vary significantly.

In the following sections, we outline the improvements in performance that `JAXtronomy` has over `lenstronomy` for each step in the pipeline. These performance benchmarks were run using an Intel(R) Xeon(R) Gold 6338 CPU @ 2.00GHz, an NVIDIA A100 GPU, and `JAX` version 0.7.0.

Deflection angle calculations

Each entry in the table indicates how much faster `JAXtronomy` is compared to `lenstronomy` at computing deflection angles for the corresponding deflector profile and grid size. Those profiles which are already computationally inexpensive for `lenstronomy` are excluded from this table. Some comparisons vary significantly with values of function arguments, so a range is given rather than a number.

Deflector Profile	60x60 grid (cpu)	180x180 grid (cpu)	180x180 grid (gpu)
CSE	1.6x	2.6x	2.6x
EPL	5.1x - 15x	9.2x - 17x	37x - 120x
EPL (jax) vs EPL_NUMBA	1.4x	3.0x	13x
EPL_MULTIPOLE_M1M3M4	2.1x - 7x	6.4x - 13x	42x - 108x
HERNQUIST	2.0x	3.4x	5.8x
HERNQUIST_ELLIPSE_CSE	3.8x	5.4x	40x
MULTIPOLE	0.9x	1.0x	8.3x - 14x
MULTIPOLE_ELL	1.5x - 2.1x	2.0x - 2.8x	70x
NFW	1.6x	3.3x	4.5x
NFW_ELLIPSE_CSE	4.1x	6.7x	31x
TNFW	2.4x	5.8x	7.5x

[3] https://github.com/google-research/google-research/tree/master/jaxstronomy

Flux calculations

An analogous table for the different light profiles is shown below.

Light Profile	60x60 grid (cpu)	180x180 grid (cpu)	180x180 grid (gpu)
CORE_SERSIC	2.0x	6.7x	4.2x
GAUSSIAN	1.0x	2.5x	1.3x
GAUSSIAN_ELLIPSE	1.5x	3.6x	2.0x
SERSIC	1.0x	1.7x	3.9x
SERSIC_ELLIPSE	1.9x	5.7x	3.2x
SHAPELETS (n_max=6)	6.2x	3.4x	15x
SHAPELETS (n_max=10)	6.0x	4.5x	17x

FFT Convolution

We find that FFT convolution using `JAX` on CPU results in variable performance boosts or slowdowns compared to `lenstronomy` (which uses `scipy`'s FFT convolution). On a 60x60 grid, and kernel sizes ranging from 3 to 45, `JAX` on CPU ranges from being 1.1x to 2.9x faster than `lenstronomy`, with no obvious correlation to kernel size. On a 180x180 grid, and kernel sizes ranging from 9 to 135, `JAXtronomy` on CPU ranges from being 0.7x to 2.5x as fast as `lenstronomy`, with no obvious correlation to kernel size.

However, FFT convolution using `JAX` on GPU is significantly faster than `scipy`. On a 60x60 grid, and kernel sizes ranging from 3 to 45, `JAX` on GPU ranges from being 1.5x to 3.5x faster than `lenstronomy`, with `JAX` performing better at higher kernel sizes. On a 180x180 grid, and kernel sizes ranging from 9 to 135, `JAXtronomy` on GPU is about 10x to 20x as fast as `lenstronomy`, again with `JAX` performing better at higher kernel sizes.

Improvements over lenstronomy in lens modelling

The process of lens modelling involves finding best-fit parameters describing a lensed system from real data. In `lenstronomy`, this typically involves a Particle Swarm Optimizer (PSO) (Kennedy & Eberhart, 1995) for optimization and Monte Carlo Markov Chains for posterior sampling.

`JAXtronomy` retains all of the lens modelling algorithms from `lenstronomy` while benefitting from the increased performance outlined above. Additionally, using `JAX`'s autodifferentiation, we have implemented the L-BFGS gradient descent algorithm from the `Optax`[4] library (DeepMind et al., 2020) for optimization. This is a significant improvement over `lenstronomy`'s PSO, which does not have access to gradient information.

References

Birrer, S. (2021). Gravitational lensing formalism in a curved arc basis: A continuous description of observables and degeneracies from the weak to the strong lensing regime. *The Astrophysical Journal*, *919*(1), 38. https://doi.org/10.3847/1538-4357/ac1108

Birrer, S., & Amara, A. (2018). Lenstronomy: Multi-purpose gravitational lens modelling software package. *Physics of the Dark Universe*, *22*, 189–201. https://doi.org/10.1016/j.dark.2018.11.002

Birrer, Simon, & Treu, Tommaso. (2021). TDCOSMO - v. Strategies for precise and accurate measurements of the hubble constant with strong lensing. *Astronomy & Astrophysics*, *649*, A61. https://doi.org/10.1051/0004-6361/202039179

[4]https://github.com/google-deepmind/optax

Birrer, S., Shajib, A. J., Galan, A., Millon, M., Treu, T., Agnello, A., Auger, M., Chen, G. C.-F., Christensen, L., Collett, T., Courbin, F., Fassnacht, C. D., Koopmans, L. V. E., Marshall, P. J., Park, J.-W., Rusu, C. E., Sluse, D., Spiniello, C., Suyu, S. H., … Van de Vyvere, L. (2020). TDCOSMO - IV. Hierarchical time-delay cosmography – joint inference of the hubble constant and galaxy density profiles⋆. *Astronomy & Astrophysics*, *643*, A165. https://doi.org/10.1051/0004-6361/202038861

Collaboration, T., Birrer, S., Buckley-Geer, E. J., Cappellari, M., Courbin, F., Dux, F., Fassnacht, C. D., Frieman, J. A., Galan, A., Gilman, D., Huang, X.-Y., Knabel, S., Langeroodi, D., Lin, H., Millon, M., Morishita, T., Motta, V., Mozumdar, P., Paic, E., … Wong, K. C. (2025). *TDCOSMO 2025: Cosmological constraints from strong lensing time delays*. https://arxiv.org/abs/2506.03023

DeepMind, Babuschkin, I., Baumli, K., Bell, A., Bhupatiraju, S., Bruce, J., Buchlovsky, P., Budden, D., Cai, T., Clark, A., Danihelka, I., Dedieu, A., Fantacci, C., Godwin, J., Jones, C., Hemsley, R., Hennigan, T., Hessel, M., Hou, S., … Viola, F. (2020). *The DeepMind JAX Ecosystem*. http://github.com/google-deepmind

Ding, X., Silverman, J., Treu, T., Schulze, A., Schramm, M., Birrer, S., Park, D., Jahnke, K., Bennert, V. N., Kartaltepe, J. S., Koekemoer, A. M., Malkan, M. A., & Sanders, D. (2020). The Mass Relations between Supermassive Black Holes and Their Host Galaxies at $1 < z < 2$ HST-WFC3. *The Astrophysical Journal*, *888*(1), 37. https://doi.org/10.3847/1538-4357/ab5b90

Galan, A., Vernardos, G., Peel, A., Courbin, F., & Starck, J.-L. (2022). Using wavelets to capture deviations from smoothness in galaxy-scale strong lenses. *Astronomy & Astrophysics*, *668*, A155. https://doi.org/10.1051/0004-6361/202244464

Gilman, D., Birrer, S., Nierenberg, A., Treu, T., Du, X., & Benson, A. (2019). Warm dark matter chills out: Constraints on the halo mass function and the free-streaming length of dark matter with eight quadruple-image strong gravitational lenses. *Monthly Notices of the Royal Astronomical Society*, *491*(4), 6077–6101. https://doi.org/10.1093/mnras/stz3480

Gilman, D., Birrer, S., & Treu, T. (2020). TDCOSMO - III. Dark matter substructure meets dark energy. The effects of (sub)halos on strong-lensing measurements of H0. *Astronomy & Astrophysics*, *642*, A194. https://doi.org/10.1051/0004-6361/202038829

Gu, A., Huang, X., Sheu, W., Aldering, G., Bolton, A. S., Boone, K., Dey, A., Filipp, A., Jullo, E., Perlmutter, S., Rubin, D., Schlafly, E. F., Schlegel, D. J., Shu, Y., & Suyu, S. H. (2022). GIGA-lens: Fast bayesian inference for strong gravitational lens modeling. *The Astrophysical Journal*, *935*(1), 49. https://doi.org/10.3847/1538-4357/ac6de4

Kennedy, J., & Eberhart, R. (1995). Particle swarm optimization. *Proceedings of ICNN'95 - International Conference on Neural Networks*, *4*, 1942–1948 vol.4. https://doi.org/10.1109/ICNN.1995.488968

Millon, M., Galan, A., Courbin, F., Treu, T., Suyu, S. H., Ding, X., Birrer, S., Chen, G. C.-F., Shajib, A. J., Sluse, D., Wong, K. C., Agnello, A., Auger, M. W., Buckley-Geer, E. J., Chan, J. H. H., Collett, T., Fassnacht, C. D., Hilbert, S., Koopmans, L. V. E., … Van de Vyvere, L. (2020). TDCOSMO - i. An exploration of systematic uncertainties in the inference of H0 from time-delay cosmography. *Astronomy & Astrophysics*, *639*, A101. https://doi.org/10.1051/0004-6361/201937351

Schmidt, T., Treu, T., Birrer, S., Millon, M., Sluse, D., Galan, A., Shajib, A., Lemon, C., Dux, F., & Courbin, F. (2025). TDCOSMO. XVIII. Strong lens model and time-delay predictions for J1721+8842, the first einstein zigzag lens. *Astronomy & Astrophysics*. https://doi.org/10.1051/0004-6361/202449984

Shajib, Anowar J., Nihal, N. S., Tan, C. Y., Sahu, V., Birrer, S., Treu, T., & Frieman, J. (2025). *dolphin: A fully automated forward modeling pipeline powered by artificial intelligence for* *galaxy-scale strong lenses*. https://arxiv.org/abs/2503.22657

Shajib, Anowar J., Treu, T., Birrer, S., & Sonnenfeld, A. (2021). Dark matter haloes of massive elliptical galaxies at $z \sim 0.2$ are well described by the Navarro-Frenk-White profile. *Monthly Notices of the Royal Astronomical Society*, *503*(2), 2380–2405. https://doi.org/10.1093/mnras/stab536

Shajib, A. J., Wong, K. C., Birrer, S., Suyu, S. H., Treu, T., Buckley-Geer, E. J., Lin, H., Rusu, C. E., Poh, J., Palmese, A., Agnello, A., Auger-Williams, M. W., Galan, A., Schuldt, S., Sluse, D., Courbin, F., Frieman, J., & Millon, M. (2022). TDCOSMO. IX. Systematic comparison between lens modelling software programs: Time-delay prediction for WGD 2038−4008. *Astronomy & Astrophysics*, *667*, A123. https://doi.org/10.1051/0004-6361/202243401

Sheu, W., Shajib, A. J., Treu, T., Sonnenfeld, A., Birrer, S., Cappellari, M., Oldham, L. J., & Tan, C. Y. (2025). Project Dinos II: redshift evolution of dark and luminous matter density profiles in strong-lensing elliptical galaxies across $0.1 < z < 0.9$. *Monthly Notices of the Royal Astronomical Society*, *541*(1), 1–27. https://doi.org/10.1093/mnras/staf976

Tan, C. Y., Shajib, A. J., Birrer, S., Sonnenfeld, A., Treu, T., Wells, P., Williams, D., Buckley-Geer, E. J., Drlica-Wagner, A., & Frieman, J. (2024). Project Dinos I: A joint lensing-dynamics constraint on the deviation from the power law in the mass profile of massive ellipticals. *Monthly Notices of the Royal Astronomical Society*, *530*(2), 1474–1505. https://doi.org/10.1093/mnras/stae884

Wagner-Carena, S., Aalbers, J., Birrer, S., Nadler, E. O., Darragh-Ford, E., Marshall, P. J., & Wechsler, R. H. (2023). From images to dark matter: End-to-end inference of substructure from hundreds of strong gravitational lenses. *The Astrophysical Journal*, *942*(2), 75. https://doi.org/10.3847/1538-4357/aca525

Wagner-Carena, S., Lee, J., Pennington, J., Aalbers, J., Birrer, S., & Wechsler, R. H. (2024). A strong gravitational lens is worth a thousand dark matter halos: Inference on small-scale structure using sequential methods. *The Astrophysical Journal*, *975*(2), 297. https://doi.org/10.3847/1538-4357/ad6e70

Wang, H., Suyu, S. H., Galan, A., Halkola, A., Cappellari, M., Shajib, A. J., & Cernetic, M. (2025). *GPU-accelerated gravitational lensing & dynamical (GLaD) modeling for cosmology and galaxies*. https://arxiv.org/abs/2504.01302

Wedig, B., Daylan, T., Birrer, S., Cyr-Racine, F.-Y., Dvorkin, C., Finkbeiner, D. P., Huang, A., Huang, X., Karthik, R., Khadka, N., Natarajan, P., Nierenberg, A. M., Peter, A. H. G., Pierel, J. D. R., Tang, X. T., & Wechsler, R. H. (2025). The roman view of strong gravitational lenses. *The Astrophysical Journal*, *986*(1), 42. https://doi.org/10.3847/1538-4357/adc24f