

# <sup>1</sup> JAXtronomy: A JAX port of lenstronomy

<sup>2</sup> Alan Huang  <sup>1</sup>, Simon Birrer  <sup>1</sup>, Natalie B. Hogg  <sup>2</sup>, Aymeric Galan  <sup>3,4</sup>,  
<sup>3</sup> Daniel Gilman  <sup>5</sup>, Anowar J. Shajib  <sup>5,6,7</sup>, and Nan Zhang  <sup>8</sup>

<sup>4</sup> 1 Department of Physics and Astronomy, Stony Brook University, Stony Brook, NY 11794, USA  
<sup>5</sup> Laboratoire Univers et Particules de Montpellier, CNRS and Université de Montpellier (UMR-5299),  
<sup>6</sup> 34095 Montpellier, France <sup>3</sup> Max-Planck-Institut für Astrophysik, Karl-Schwarzschild Straße 1, 85748  
<sup>7</sup> Garching, Germany <sup>4</sup> Technical University of Munich, TUM School of Natural Sciences, Physics  
<sup>8</sup> Department, James-Franck-Straße 1, 85748 Garching, Germany <sup>5</sup> Department of Astronomy and  
<sup>9</sup> Astrophysics, University of Chicago, Chicago, Illinois 60637, USA <sup>6</sup> Kavli Institute for Cosmological  
<sup>10</sup> Physics, University of Chicago, Chicago, IL 60637, USA <sup>7</sup> Center for Astronomy, Space Science and  
<sup>11</sup> Astrophysics, Independent University, Bangladesh, Dhaka 1229, Bangladesh <sup>8</sup> Department of Physics,  
<sup>12</sup> University of Illinois, 1110 West Green St., Urbana, IL 61801, USA

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

## Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

## License

Authors of papers retain copyright <sup>23</sup> and release the work under a <sup>24</sup> Creative Commons Attribution 4.0 <sup>25</sup> International License ([CC BY 4.0](#)). <sup>26</sup> <sup>27</sup>

## <sup>13</sup> Summary

<sup>14</sup> Gravitational lensing is a phenomenon where light bends around massive objects, resulting in <sup>15</sup> distorted images seen by an observer. Studying gravitationally lensed objects can give us key <sup>16</sup> insights into cosmology and astrophysics, such as constraints on the expansion rate of the <sup>17</sup> universe and dark matter models.

<sup>18</sup> Thus, we introduce JAXtronomy, a re-implementation of the gravitational lensing software <sup>19</sup> package `lenstronomy`<sup>1</sup> ([Birrer, 2021; Birrer & Amara, 2018](#)) using JAX<sup>2</sup>. JAX is a Python <sup>20</sup> library that uses an accelerated linear algebra (XLA) compiler to improve the performance of <sup>21</sup> computing software. Our core design principle of JAXtronomy is to maintain an identical API <sup>22</sup> to that of `lenstronomy`.

<sup>23</sup> The main JAX features utilized in JAXtronomy are just-in-time-compilation, which can lead to <sup>24</sup> significant reductions in execution time, and automatic differentiation, which allows for the <sup>25</sup> implementation of gradient-based algorithms that were previously impossible. Additionally, <sup>26</sup> JAX allows code to be run on GPUs or parallelized across CPU cores, further boosting the <sup>27</sup> performance of JAXtronomy.

## <sup>28</sup> Statement of need

<sup>29</sup> `lenstronomy` has been widely applied to numerous science cases, with more than 200 <sup>30</sup> publications making use of the software, and has an increasing number of dependent packages <sup>31</sup> relying on features of `lenstronomy`. For instance, science cases directly involving `lenstronomy` <sup>32</sup> include galaxy evolution studies using strong lensing ([Anowar J. Shajib et al., 2021; Sheu et al., 2025; Tan et al., 2024](#)) and detailed lens modeling for measuring the Hubble constant <sup>33</sup> using time-delay cosmography by the TDCOSMO collaboration ([Birrer, S. et al., 2020; Birrer, Simon & Treu, Tommaso, 2021; Collaboration et al., 2025; Gilman, D. et al., 2020; Millon, M. et al., 2020; Schmidt et al., 2025; A. J. Shajib et al., 2022; ?](#)).

<sup>37</sup> Examples of packages dependent on `lenstronomy` for general-purpose lensing computations and <sup>38</sup> image modelling include the `dolphin` package ([Anowar J. Shajib et al., 2025](#)) for automated <sup>39</sup> lens modeling, the `galight` package ([Ding et al., 2020](#)) for galaxy morphology measurements,

<sup>1</sup><https://github.com/lenstronomy/lenstronomy>

<sup>2</sup><https://github.com/jax-ml/jax>

<sup>40</sup> SLSim (Khadka et al, 2025, in prep) for simulating large populations of strong lenses, pyHalo  
<sup>41</sup> (Gilman et al., 2019) and mejiro (Wedig et al., 2025) for simulating strong lenses with dark  
<sup>42</sup> matter substructure, and PALTAS (Wagner-Carena et al., 2023) for neural network inference  
<sup>43</sup> tasks.

<sup>44</sup> In many of these applications, computational constraints are the key limiting factor for strong  
<sup>45</sup> gravitational lensing science. For example, increased data quality and number of lenses to  
<sup>46</sup> analyze makes lens modeling a computational bottleneck, and expensive ray-tracing through  
<sup>47</sup> tens of thousands of dark matter substructures limit the amount of images that can be  
<sup>48</sup> simulated, especially for the training of neural networks and simulation-based inferences. These  
<sup>49</sup> ever-increasing computational costs have lead to the development of several JAX-accelerated  
<sup>50</sup> and GPU-accelerated strong-lensing packages, such as gigalens (Gu et al., 2022), herculens  
<sup>51</sup> (Galan et al., 2022), paltax (Wagner-Carena et al., 2024), GLaD (Wang et al., 2025), caustics<sup>3</sup>  
<sup>52</sup> (Stone et al., 2024) and Google Research's jaxstronomy<sup>4</sup>.

### <sup>53</sup> Why JAXtronomy?

<sup>54</sup> JAXtronomy inherits a wide range of features from lenstronomy that are not offered by any  
<sup>55</sup> of the aforementioned JAX-accelerated or GPU-accelerated software. Such features include  
<sup>56</sup> lenstronomy's linear amplitude solver, which reduces the number of sampled parameters during  
<sup>57</sup> lens modeling, as well as a variety of log likelihood functions and optional punishment terms.  
<sup>58</sup> Furthermore, JAXtronomy aims to maintain an identical API to lenstronomy so that packages  
<sup>59</sup> dependent on lenstronomy can transition seamlessly to JAXtronomy.

## <sup>60</sup> Improvements over lenstronomy in image simulation

<sup>61</sup> The simulation of a lensed image comes in three main steps. The first step begins with a  
<sup>62</sup> coordinate grid in the angles seen by the observer. These coordinates are ray-traced through  
<sup>63</sup> the deflectors back to the source plane. This process requires the calculation of light ray  
<sup>64</sup> deflection angles at each deflector. Second, the surface brightness of the source is calculated  
<sup>65</sup> on the ray-traced coordinate grid. This produces a lensed image. Third, the lensed image gets  
<sup>66</sup> convolved by the point spread function (PSF) originating from diffraction of the telescope  
<sup>67</sup> optics and atmospheric turbulence. Due to the various choices in deflector mass profiles, light  
<sup>68</sup> model profiles, grid size, and PSF kernel size, the overall runtime of the pipeline can vary  
<sup>69</sup> significantly.

<sup>70</sup> In the following sections, we outline the improvements in performance that JAXtronomy has  
<sup>71</sup> over lenstronomy for each step in the pipeline. These performance benchmarks were run using  
<sup>72</sup> an Intel(R) Xeon(R) Gold 6338 CPU @ 2.00GHz, an NVIDIA A100 GPU, and JAX version  
<sup>73</sup> 0.7.0.

### <sup>74</sup> Deflection angle calculations

<sup>75</sup> Each entry in the table indicates how much faster JAXtronomy is compared to lenstronomy  
<sup>76</sup> at computing deflection angles for the corresponding deflector profile and grid size. Some  
<sup>77</sup> comparisons vary significantly with values of function arguments, so a range is given rather  
<sup>78</sup> than a number.

Deflector Profile	60x60 grid (cpu)	180x180 grid (cpu)	180x180 grid (gpu)
CONVERGENCE	0.4x	1.1x	0.5x
CSE	1.6x	2.6x	2.6x
EPL	5.1x - 15x	9.2x - 17x	37x - 120x

<sup>3</sup><https://github.com/Ciela-Institute/caustics>

<sup>4</sup><https://github.com/google-research/google-research/tree/master/jaxstronomy>

Deflector Profile	60x60 grid (cpu)	180x180 grid (cpu)	180x180 grid (gpu)
EPL (jax) vs EPL_NUMBA	1.4x	3.0x	13x
EPL_MULTIPOLE_M1M3M4	2.1x - 7x	6.4x - 13x	42x - 108x
HERNQUIST	2.0x	3.4x	5.8x
HERNQUIST_ELLIPSE_CSE	3.8x	5.4x	40x
MULTIPOLE	0.9x	1.0x	8.3x - 14x
MULTIPOLE_ELL	1.5x - 2.1x	2.0x - 2.8x	70x
NIE/SIE	0.5x	0.5x	2.0x
NFW	1.6x	3.3x	4.5x
NFW_ELLIPSE_CSE	4.1x	6.7x	31x
PJAFFE	1.0x	1.2x	2.8x
PJAFFE_ELLIPSE_POTENTIAL	1.4x	1.6x	3.1x
SHEAR	0.7x	2.0x	0.9x
SIS	1.4x	3.3x	2.0x
TNFW	2.4x	5.8x	7.5x

<sup>79</sup> For small enough grid sizes, JAXtronomy computes deflection angles slower than lenstronomy  
<sup>80</sup> when using certain deflector profiles. This is because function call overheads are significantly  
<sup>81</sup> higher in JAX than in standard Python, so computations that are already fast in Python  
<sup>82</sup> can end up slower in JAX. In these cases, the benefit of using JAX is to have automatic  
<sup>83</sup> differentiation for lens modeling.

### <sup>84</sup> Flux calculations

<sup>85</sup> An analogous table for the different light profiles is shown below. The MULTI\_GAUSSIAN  
<sup>86</sup> and MULTI\_GAUSSIAN\_ELLIPSE profiles include five GAUSSIAN and GAUSSIAN\_ELLIPSE  
<sup>87</sup> components, respectively, highlighting JAX's improved performance in sequential computations.

Light Profile	60x60 grid (cpu)	180x180 grid (cpu)	180x180 grid (gpu)
CORE_SERSIC	2.0x	6.7x	4.2x
GAUSSIAN	1.0x	2.5x	1.3x
GAUSSIAN_ELLIPSE	1.5x	3.6x	2.0x
MULTI_GAUSSIAN	3.7x	11x	7.8x
MULTI_GAUSSIAN_ELLIPSE	4.0x	13x	6.9x
SERSIC	1.0x	1.7x	3.9x
SERSIC_ELLIPSE	1.9x	5.7x	3.2x
SERSIC_ELLIPSE_Q_PHI	1.7x	5.5x	3.3x
SHAPELETS (n_max=6)	6.2x	3.4x	15x
SHAPELETS (n_max=10)	6.0x	4.5x	17x

### <sup>88</sup> FFT Convolution

<sup>89</sup> We find that FFT convolution using JAX on CPU results in variable performance boosts or  
<sup>90</sup> slowdowns compared to lenstronomy (which uses scipy's FFT convolution). On a 60x60  
<sup>91</sup> grid, and kernel sizes ranging from 3 to 45, JAX on CPU ranges from being 1.1x to 2.9x faster  
<sup>92</sup> than lenstronomy, with no obvious correlation to kernel size. On a 180x180 grid, and kernel  
<sup>93</sup> sizes ranging from 9 to 135, JAXtronomy on CPU ranges from being 0.7x to 2.5x as fast as  
<sup>94</sup> lenstronomy, with no obvious correlation to kernel size.

95     However, FFT convolution using JAX on GPU is significantly faster than `scipy`. On a 60x60  
 96     grid, and kernel sizes ranging from 3 to 45, JAX on GPU ranges from being 1.5x to 3.5x  
 97     faster than `lenstronomy`, with JAX performing better at higher kernel sizes. On a 180x180  
 98     grid, and kernel sizes ranging from 9 to 135, `JAXtronomy` on GPU is about 10x to 20x as fast  
 99     as `lenstronomy`, again with JAX performing better at higher kernel sizes.

## 100    Improvements over `lenstronomy` in lens modelling

101    The process of lens modelling involves finding best-fit parameters describing a lensed system from  
 102    real data. In `lenstronomy`, this typically involves a Particle Swarm Optimizer (PSO) ([Kennedy & Eberhart, 1995](#)) for optimization and Monte Carlo Markov Chains for posterior sampling.  
 103    `JAXtronomy` retains these lens modelling algorithms from `lenstronomy` while benefitting from  
 104    the increased performance outlined above.

105    In the following table, we compare `JAXtronomy`'s PSO performance to that of `lenstronomy`  
 106    when modeling a lens with an elliptical power-law mass profile, Sersic light profile, and a  
 107    quadruply-imaged point source. We use a 100x100 grid and a size 13 PSF kernel. These  
 108    benchmarks were performed using the same hardware as in the previous section.

Number of Particles	1 CPU core	64 CPU cores (parallelized)	GPU
64	4x	16x	5x
128	4x	18x	5.5x
256	4.7x	30x	9x
512	4.7x	34x	11x

110    Additionally, using JAX's autodifferentiation, we have implemented the L-BFGS gradient descent  
 111    algorithm from the Optax<sup>5</sup> library ([DeepMind et al., 2020](#)) for optimization. This is a significant  
 112    improvement over `lenstronomy`'s PSO, which does not have access to gradient information.  
 113    Due to the random nature of the PSO, we do not present a concrete comparison between  
 114    `lenstronomy` and `JAXtronomy` for how long it takes to find best-fit parameters. However, we  
 115    note that `JAXtronomy` can find a good fit within one minute, while `lenstronomy` can take  
 116    hours.

## 117    References

- 118    Birrer, S. (2021). Gravitational lensing formalism in a curved arc basis: A continuous  
 119    description of observables and degeneracies from the weak to the strong lensing regime.  
 120    *The Astrophysical Journal*, 919(1), 38. <https://doi.org/10.3847/1538-4357/ac1108>
- 121    Birrer, S., & Amara, A. (2018). Lenstronomy: Multi-purpose gravitational lens modelling  
 122    software package. *Physics of the Dark Universe*, 22, 189–201. <https://doi.org/10.1016/j.dark.2018.11.002>
- 123    Birrer, Simon, & Treu, Tommaso. (2021). TDCOSMO - v. Strategies for precise and accurate  
 124    measurements of the hubble constant with strong lensing. *Astronomy & Astrophysics*, 649,  
 125    A61. <https://doi.org/10.1051/0004-6361/202039179>
- 126    Birrer, S., Shajib, A. J., Galan, A., Millon, M., Treu, T., Agnello, A., Auger, M., Chen, G.  
 127    C.-F., Christensen, L., Collett, T., Courbin, F., Fassnacht, C. D., Koopmans, L. V. E.,  
 128    Marshall, P. J., Park, J.-W., Rusu, C. E., Sluse, D., Spinelli, C., Suyu, S. H., ... Van de  
 129    Vyvere, L. (2020). TDCOSMO - IV. Hierarchical time-delay cosmography – joint inference  
 130    of the hubble constant and galaxy density profiles\*. *Astronomy & Astrophysics*, 643, A165.  
 131    <https://doi.org/10.1051/0004-6361/202038861>

<sup>5</sup><https://github.com/google-deepmind/optax>

- 133 Collaboration, T., Birrer, S., Buckley-Geer, E. J., Cappellari, M., Courbin, F., Dux, F.,  
 134 Fassnacht, C. D., Frieman, J. A., Galan, A., Gilman, D., Huang, X.-Y., Knobel, S.,  
 135 Langeroodi, D., Lin, H., Millon, M., Morishita, T., Motta, V., Mozumdar, P., Paic, E., ...  
 136 Wong, K. C. (2025). *TDCOSMO 2025: Cosmological constraints from strong lensing time*  
 137 *delays*. <https://arxiv.org/abs/2506.03023>
- 138 DeepMind, Babuschkin, I., Baumli, K., Bell, A., Bhupatiraju, S., Bruce, J., Buchlovsky, P.,  
 139 Budden, D., Cai, T., Clark, A., Danihelka, I., Dedieu, A., Fantacci, C., Godwin, J., Jones,  
 140 C., Hemsley, R., Hennigan, T., Hessel, M., Hou, S., ... Viola, F. (2020). *The DeepMind*  
 141 *JAX Ecosystem*. <http://github.com/google-deepmind>
- 142 Ding, X., Silverman, J., Treu, T., Schulze, A., Schramm, M., Birrer, S., Park, D., Jahnke, K.,  
 143 Bennert, V. N., Kartaltepe, J. S., Koekemoer, A. M., Malkan, M. A., & Sanders, D. (2020).  
 144 The Mass Relations between Supermassive Black Holes and Their Host Galaxies at  $1 < z$   
 145  $< 2$  HST-WFC3. *The Astrophysical Journal*, 888(1), 37. <https://doi.org/10.3847/1538-4357/ab5b90>
- 147 Galan, A., Vernardos, G., Peel, A., Courbin, F., & Starck, J.-L. (2022). Using wavelets  
 148 to capture deviations from smoothness in galaxy-scale strong lenses. *Astronomy &*  
 149 *Astrophysics*, 668, A155. <https://doi.org/10.1051/0004-6361/202244464>
- 150 Gilman, D., Birrer, S., Nierenberg, A., Treu, T., Du, X., & Benson, A. (2019). Warm dark  
 151 matter chills out: Constraints on the halo mass function and the free-streaming length of  
 152 dark matter with eight quadruple-image strong gravitational lenses. *Monthly Notices of the*  
 153 *Royal Astronomical Society*, 491(4), 6077–6101. <https://doi.org/10.1093/mnras/stz3480>
- 154 Gilman, D., Birrer, S., & Treu, T. (2020). TDCOSMO - III. Dark matter substructure meets  
 155 dark energy. The effects of (sub)halos on strong-lensing measurements of  $H_0$ . *Astronomy*  
 156 & *Astrophysics*, 642, A194. <https://doi.org/10.1051/0004-6361/202038829>
- 157 Gu, A., Huang, X., Sheu, W., Aldering, G., Bolton, A. S., Boone, K., Dey, A., Filipp, A.,  
 158 Jullo, E., Perlmutter, S., Rubin, D., Schlafly, E. F., Schlegel, D. J., Shu, Y., & Suyu, S. H.  
 159 (2022). GIGA-lens: Fast bayesian inference for strong gravitational lens modeling. *The*  
 160 *Astrophysical Journal*, 935(1), 49. <https://doi.org/10.3847/1538-4357/ac6de4>
- 161 Kennedy, J., & Eberhart, R. (1995). Particle swarm optimization. *Proceedings of ICNN'95*  
 162 - *International Conference on Neural Networks*, 4, 1942–1948 vol.4. <https://doi.org/10.1109/ICNN.1995.488968>
- 164 Millon, M., Galan, A., Courbin, F., Treu, T., Suyu, S. H., Ding, X., Birrer, S., Chen, G. C.-F.,  
 165 Shajib, A. J., Sluse, D., Wong, K. C., Agnello, A., Auger, M. W., Buckley-Geer, E. J.,  
 166 Chan, J. H. H., Collett, T., Fassnacht, C. D., Hilbert, S., Koopmans, L. V. E., ... Van  
 167 de Vyvere, L. (2020). TDCOSMO - i. An exploration of systematic uncertainties in the  
 168 inference of  $H_0$  from time-delay cosmography. *Astronomy & Astrophysics*, 639, A101.  
 169 <https://doi.org/10.1051/0004-6361/201937351>
- 170 Schmidt, T., Treu, T., Birrer, S., Millon, M., Sluse, D., Galan, A., Shajib, A., Lemon, C.,  
 171 Dux, F., & Courbin, F. (2025). TDCOSMO. XVIII. Strong lens model and time-delay  
 172 predictions for J1721+8842, the first einstein zigzag lens. *Astronomy & Astrophysics*.  
 173 <https://doi.org/10.1051/0004-6361/202449984>
- 174 Shajib, Anowar J., Nihal, N. S., Tan, C. Y., Sahu, V., Birrer, S., Treu, T., & Frieman, J. (2025).  
 175 *dolphin: A fully automated forward modeling pipeline powered by artificial intelligence for*  
 176 *galaxy-scale strong lenses*. <https://arxiv.org/abs/2503.22657>
- 177 Shajib, Anowar J., Treu, T., Birrer, S., & Sonnenfeld, A. (2021). Dark matter haloes of  
 178 massive elliptical galaxies at  $z \sim 0.2$  are well described by the Navarro-Frenk-White  
 179 profile. *Monthly Notices of the Royal Astronomical Society*, 503(2), 2380–2405. <https://doi.org/10.1093/mnras/stab536>
- 181 Shajib, A. J., Wong, K. C., Birrer, S., Suyu, S. H., Treu, T., Buckley-Geer, E. J., Lin, H., Rusu, C.

- 182 E., Poh, J., Palmese, A., Agnello, A., Auger-Williams, M. W., Galan, A., Schuldt, S., Sluse,  
183 D., Courbin, F., Frieman, J., & Millon, M. (2022). TDCOSMO. IX. Systematic comparison  
184 between lens modelling software programs: Time-delay prediction for WGD 2038—4008.  
185 *Astronomy & Astrophysics*, 667, A123. <https://doi.org/10.1051/0004-6361/202243401>
- 186 Sheu, W., Shajib, A. J., Treu, T., Sonnenfeld, A., Birrer, S., Cappellari, M., Oldham, L. J., &  
187 Tan, C. Y. (2025). Project Dinos II: redshift evolution of dark and luminous matter density  
188 profiles in strong-lensing elliptical galaxies across  $0.1 < z < 0.9$ . *Monthly Notices of the*  
189 *Royal Astronomical Society*, 541(1), 1–27. <https://doi.org/10.1093/mnras/staf976>
- 190 Stone, C., Adam, A., Coogan, A., Yantovski-Barth, M. J., Filipp, A., Setiawan, L., Core,  
191 C., Legin, R., Wilson, C., Barco, G. M., Hezaveh, Y., & Perreault-Levasseur, L. (2024).  
192 Caustics: A python package for accelerated strong gravitational lensing simulations. *Journal*  
193 *of Open Source Software*, 9(103), 7081. <https://doi.org/10.21105/joss.07081>
- 194 Tan, C. Y., Shajib, A. J., Birrer, S., Sonnenfeld, A., Treu, T., Wells, P., Williams, D.,  
195 Buckley-Geer, E. J., Drlica-Wagner, A., & Frieman, J. (2024). Project Dinos I: A joint  
196 lensing-dynamics constraint on the deviation from the power law in the mass profile of  
197 massive ellipticals. *Monthly Notices of the Royal Astronomical Society*, 530(2), 1474–1505.  
198 <https://doi.org/10.1093/mnras/stae884>
- 199 Wagner-Carena, S., Aalbers, J., Birrer, S., Nadler, E. O., Darragh-Ford, E., Marshall, P. J., &  
200 Wechsler, R. H. (2023). From images to dark matter: End-to-end inference of substructure  
201 from hundreds of strong gravitational lenses. *The Astrophysical Journal*, 942(2), 75.  
202 <https://doi.org/10.3847/1538-4357/aca525>
- 203 Wagner-Carena, S., Lee, J., Pennington, J., Aalbers, J., Birrer, S., & Wechsler, R. H.  
204 (2024). A strong gravitational lens is worth a thousand dark matter halos: Inference on  
205 small-scale structure using sequential methods. *The Astrophysical Journal*, 975(2), 297.  
206 <https://doi.org/10.3847/1538-4357/ad6e70>
- 207 Wang, H., Suyu, S. H., Galan, A., Halkola, A., Cappellari, M., Shajib, A. J., & Cernetic, M.  
208 (2025). GPU-accelerated gravitational lensing & dynamical (GLaD) modeling for cosmology  
209 and galaxies. <https://arxiv.org/abs/2504.01302>
- 210 Wedig, B., Daylan, T., Birrer, S., Cyr-Racine, F.-Y., Dvorkin, C., Finkbeiner, D. P., Huang, A.,  
211 Huang, X., Karthik, R., Khadka, N., Natarajan, P., Nierenberg, A. M., Peter, A. H. G., Pierel,  
212 J. D. R., Tang, X. T., & Wechsler, R. H. (2025). The roman view of strong gravitational  
213 lenses. *The Astrophysical Journal*, 986(1), 42. <https://doi.org/10.3847/1538-4357/adc24f>